

CS4725/CS6705

Chapter 5: Adversarial Search

Adversarial search = games

- Multi-agent environments in which the agents' goals are in conflict
- Normally focused on deterministic, turn-taking, two-player, **zero-sum** games of **perfect information**
 - e.g., one player wins (+1) and one player loses (-1)

Example: chess


- Branching factor of about 35
- Games often last about 100 moves
- Search tree has 35^{100} (or 10^{154}) nodes
(although the search graph has only 10^{40} distinct nodes)
- Calculating optimal decisions is infeasible

Optimal decisions in games


- Consider a two-player game, where MAX moves first, then MIN, then MAX, etc.
- Game definition:
 - **Initial state:** board position, whose move it is
 - **Successor function:** returns a list of (move,state) pairs for the current state
 - **Terminal test:** determines if the game is over
 - **Utility function:** provides a numerical value for terminal states (e.g., +1 for a win, -1 for a loss, 0 for a draw)
- Game tree (tic-tac-toe example on next slides)

Tic-tac-toe example from Russell&Norvig


MAX (x)




Tic-tac-toe example from Russell&Norvig




Tic-tac-toe example from Russell&Norvig



Tic-tac-toe example from Russell&Norvig



Tic-tac-toe example from Russell&Norvig



Optimal strategies

- We can't just search for a sequence of moves that will lead us to a good solution.
- We must consider the fact that, in between our moves, the *opponent* will be making moves, working *against* our goal.

Minimax value

- Minimax value of a node: the utility (for MAX) of being in the current state, *assuming that both players play optimally* for the rest of the game
- Minimax value of node n =
 - Utility(n) if n is a terminal state
 - Max. minimax value of n 's successors if n is a MAX node
 - Min. minimax value of n 's successors if n is a MIN node
- [Simple game example on board]

Minimax algorithm

- Algorithm used to compute the minimax value of the current state, based on the definition on the previous slide
- Start from the leaves of the tree and back up
- Completely impractical for anything but the simplest games, but an important basis for the development of more practical algorithms

Alpha-beta pruning

- Minimax search: number of states is exponential in the number of moves
- Alpha-beta pruning: [simple example on board]
- In this example, we just saved the work of looking at 1 leaf node; in general, we can prune entire subtrees and save a significant amount of work

Alpha-beta pruning

- General idea:
 - Consider a node n in the tree that the player might move to.
 - If a player has a better choice m either at the parent node of n or at any choice point further up, then the player will never choose to go to node n .
 - [Diagram on board]

Alpha-beta pruning

- [Depth-first search] Consider all nodes along the current path in the tree
- Definitions:
 - α : the value of the best (highest) choice we have found so far for MAX at any choice point along the path
 - β : the value of the best (lowest) choice we have found so far for MIN at any choice point along the path

Alpha-beta pruning

- Some notation: For a given *state*,
 - Let $\text{ACTIONS}(state)$ denote the set of actions available in that state.
 - Let $\text{RESULT}(state, action)$ be the state that results from performing *action* in *state*.
 - $\text{TERMINAL-TEST}(state)$ is true if *state* is a terminal state (has no successors / represents the end of the game).

Alpha-beta pruning (continued)

- Perform the following algorithm:

function ALPHA-BETA-SEARCH(*state*) **returns** an action

$v \leftarrow \text{MAX-VALUE}(\text{state}, -\infty, +\infty)$

return the *action* in $\text{ACTIONS}(\text{state})$ that has value v

(See the next two slides for MAX-VALUE and MIN-VALUE pseudocode.)

Alpha-beta pruning (continued)

```
function MAX-VALUE(state,  $\alpha$ ,  $\beta$ ) returns a utility value  
  if TERMINAL-TEST(state) then return UTILITY(state)  
   $v \leftarrow -\infty$   
  for each action in ACTIONS(state) do  
     $v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(\text{state}, \text{action}), \alpha, \beta))$   
    if  $v \geq \beta$  then return  $v$            // ignore/prune all other actions  
     $\alpha \leftarrow \text{MAX}(\alpha, v)$   
return  $v$ 
```

Alpha-beta pruning (continued)

```
function MIN-VALUE(state,  $\alpha$ ,  $\beta$ ) returns a utility value  
  if TERMINAL-TEST(state) then return UTILITY(state)  
   $v \leftarrow +\infty$   
  for each action in ACTIONS(state) do  
     $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{RESULT}(\textit{state}, \textit{action}), \alpha, \beta))$   
    if  $v \leq \alpha$  then return  $v$            // ignore/prune all other actions  
     $\beta \leftarrow \text{MIN}(\beta, v)$   
return  $v$ 
```

Alpha-beta pruning (continued)

- [Detailed example on board]

Effectiveness of alpha-beta pruning

- Depends on the order in which successors are checked
- If we do a good job of ordering successors:
 - Only about $O(b^{d/2})$ nodes have to be checked instead of $O(b^d)$
 - Therefore, we can look ahead about twice as far in the same amount of time
 - Example: effective branching factor for chess becomes 6 instead of 35

Imperfect, real-time decisions

- Even with alpha-beta pruning, the search space for most real games is far too large to find perfect game strategies.
- One approach: cut off the search tree at a reasonable depth and apply an **evaluation function** to leaf nodes:
 - An estimate of how “good” each state is
 - Chess example

Games involving chance

- Game tree must be adjusted to include chance nodes
- [Backgammon example on board]
- Instead of finding definite minimax values for nodes, we have to calculate the **expected value** over all possible results of the chance event (e.g., rolling dice)
- **expectiminimax value**

Expectiminimax value

- Expectiminimax value of node n =
 - $Utility(n)$ if n is a terminal state
 - Max. expectiminimax value of n 's successors if n is a MAX node
 - Min. expectiminimax value of n 's successors if n is a MIN node
 - If n is a chance node... sum (over all successors s of n) of:
 $Probability(s) \times expectiminimax(s)$
- [Simple example on board]