

CS4725/CS6705

Chapter 17: Sequential Decision Problems

Markov Decision Processes

- A formal model used to handle *sequential decision problems*:
 - Problems where the agent's eventual utility depends on a sequence of decisions made by the agent
- Interested in computing an *optimal policy*, which will tell the agent the best decision to make in every possible situation it could find itself in

MDP model

- A Markov Decision Process (MDP) consists of:
 - a set S of possible world states
(one state s_0 in S is the *initial state*)
 - $Actions(s)$: a set of possible actions for each state s
 - a transition model $P(s' \mid s, a)$ describing the effects of each action in each state (see next slide)
 - a reward function R (see upcoming slide)

Transition model

- For each state s and each action a , $P(s' \mid s, a)$ is the probability of reaching a next state s' if we perform action a in state s .
- Deterministic actions:
 - There will be only one next state, and its probability will be 1.
- In general:
 - There could be multiple next states, with each probability being in the range $[0,1]$.

Reward function

- Two common versions of the reward function:
 - Each time you visit a state, there is an associated reward. (*We will use this model.*)

$$R: S \rightarrow \mathbb{R}$$

- Each time you take a particular action in a particular state, there is an associated reward.

$$R: S \times A \rightarrow \mathbb{R}$$

Markov property

- We assume that the problem satisfies the *Markov property*:
 - The effects of an action depend only on the current state and not on previous states the agent has been in.

Full observability

- For now, we assume *full observability*:
 - Once an agent performs an action, it can observe what new state resulted from the action.
- Later: *partially observable* MDPs
 - The agent doesn't *know* its current state.
 - It relies on a probability distribution over states it *thinks* it might be in.

Policies

- Based on our formalization, we can now define a *policy* π :
 - For every state, what action is taken in that state?
- To execute a policy, an agent simply determines the current state s and then chooses the action $\pi(s)$ specified by the policy.

Policies

- The optimal policy is computed by determining, for each state, the action that will lead to the highest expected utility.
- For now: the utility of a sequence of states is simply additive:

$$U([s_0, s_1, s_2, \dots]) = R(s_0) + R(s_1) + R(s_2) + \dots$$

Grid-world example

- [Details on blackboard]

Some MDP terminology

- Finite horizon vs. infinite horizon
 - finite horizon: there is a fixed time after which no actions matter
 - infinite horizon: no fixed deadline
- Stationary vs. nonstationary policies
 - nonstationary: the optimal action for a given state might change, depending on current time [Grid-world example]
 - stationary: optimal action depends only on state
- We will focus on infinite-horizon problems.

Utility of state sequences

- Earlier, we assumed that the utility of a sequence of states was calculated simply by adding the rewards associated with the individual states:

$$U_h([s_0, s_1, s_2, \dots]) = R(s_0) + R(s_1) + R(s_2) + \dots$$

- There is another possibility...

Discounted rewards

- With discounted rewards, the utility of a state sequence is:

$$U_h([s_0, s_1, s_2, \dots]) = R(s_0) + \gamma R(s_1) + \gamma^2 R(s_2) + \dots$$

- where γ is a discount factor between 0 and 1.
 - γ indicates preference for current rewards over future ones
 - When $\gamma = 1$, discounted rewards = additive rewards

Optimal policy with discounted rewards

- If we are using discounted rewards, the value of a policy π is the *expected* sum of discounted rewards, taken across *all* possible sequences of states that could result from following π .
- The value of policy π is $E \left[\sum_{t=0}^{\infty} \gamma^t R(s_t) \right]$
- The optimal policy π^* is the policy with maximum value.

Computing optimal policies

- We will talk about two algorithms used for computing optimal policies:
 - Value iteration
 - Policy iteration

Value iteration

- Basic idea:
 - calculate the utility of each *state*
 - for each state, choose the action that maximizes the expected utility of the next state
- How do we compute the utility of a state...?

Bellman Equations

- The utility of a state is:
 - the immediate reward for that state
 - +
 - the expected discounted utility of the next state, assuming that the agent chooses the optimal action

$$U_{i+1}(s) = R(s) + \gamma \max_{a \in \text{Actions}(s)} \sum_{s'} P(s' | s, a) U_i(s')$$

Value iteration algorithm

- Initialize utilities of all states to arbitrary values (*e.g.*, all zero).
- Iteratively update the utility of each state (using steps called **Bellman updates**) until we reach an equilibrium.
- The algorithm is guaranteed to converge to unique solutions.
- [details on blackboard]

Policy iteration

- Start with some initial policy π_0
- Alternate between the following two steps until step 2 yields no change in the policy:

(1) policy evaluation

(2) policy improvement

[details + example on blackboard]

Modified policy iteration

- For large state spaces, solving linear equations exactly – with running time $O(n^3)$ – might be too slow.
- Alternative: Instead of doing exact policy iteration, do a sequence of simplified value iteration steps to give an approximation of the utilities. [Details omitted]

When to use value/policy iteration

- [Source: MDP tutorial by Andrew Moore]
- Lots of actions? *Policy iteration*
- Already have a decent policy? *Policy iteration*
- Few actions, acyclic? *Value iteration*
- *Modified policy iteration*: sort of the “best of both worlds”

Partially Observable MDPs (POMDPs)

- As discussed earlier, sometimes an agent will not *know* its current state. It will only have probabilistic information about the states it *might* be in.
- More difficult than fully observable MDPs, but a realistic model for real-world situations

POMDPs (formal definition)

- Same as a regular MDP, plus:
 - A sensor model $P(e|s)$: the probability of perceiving evidence e in state s
 - Example: In the grid world, suppose we know that there is a sound coming from square (4,C) that we can hear if we are 2 steps or fewer from (4,C). If we can't hear the sound, we must be in (1,A), (1,B), (1,C), (2,A) or (3,A) – with equal probability if we have no further evidence

Belief states

- A belief state b is a probability distribution over all possible states.
- Example from previous slide:

$$b = \langle 0.2, 0.2, 0.2, 0, 0.2, 0, 0, 0.2, 0, 0, 0 \rangle$$

Grid-world example

- Suppose we have no sensors at all in the grid world. Initially, we could be in *any* of the 9 nonterminal states (with equal probability).
- What could we do to ensure a pretty high probability of ending up in the good terminal state (state (4,C), with reward +1)?
- [details on blackboard]

Belief states

- Current belief state b' calculated as the conditional probability distribution over all states, using the previous belief state b and the new observation

$$b'(s') = \alpha \cdot P(e | s') \sum_s P(s' | s, a) b(s)$$

where α is a normalizing constant that makes the belief state sum to 1.

- Belief state calculation example [on blackboard]

Policies in POMDPs

- Key fact: The optimal action in a POMDP depends on the agent's current *belief state*, not on the actual state it is in.
- Optimal policy is a mapping $\pi^*(b)$ from belief states to actions

Decisions in POMDPs

- Decision cycle:
 - Given the current belief state b , execute the action $a = \pi^*(b)$.
 - Receive percept e .
 - Compute the new belief state b' .

POMDPs vs. MDPs

- *Note:* Solving a POMDP can be reduced to solving a fully observable MDP on the corresponding belief state space.
- However, this new MDP involves a continuous state space and standard value iteration and policy iteration algorithms will not work.
- [More complex algorithms, beyond the scope of this course, are needed.]

Grid-world example revisited

- Recall our grid-world example where we had no idea where we were initially (equal probability for each of the 9 nonterminal states).
- We came up with an idea for a policy:
Left x 5, Up x 5, Right x 5
- 77.5% probability of reaching the +1 state, expected utility of 0.08

Grid-world example revisited

- Optimal policy turns out to be:

L, U, U, R, U, U, R, U, U, R, U, R, U, R, U, ...

- 86.6% probability of reaching the +1 state,
expected utility of 0.38